



JDBC and Database Programming in Java

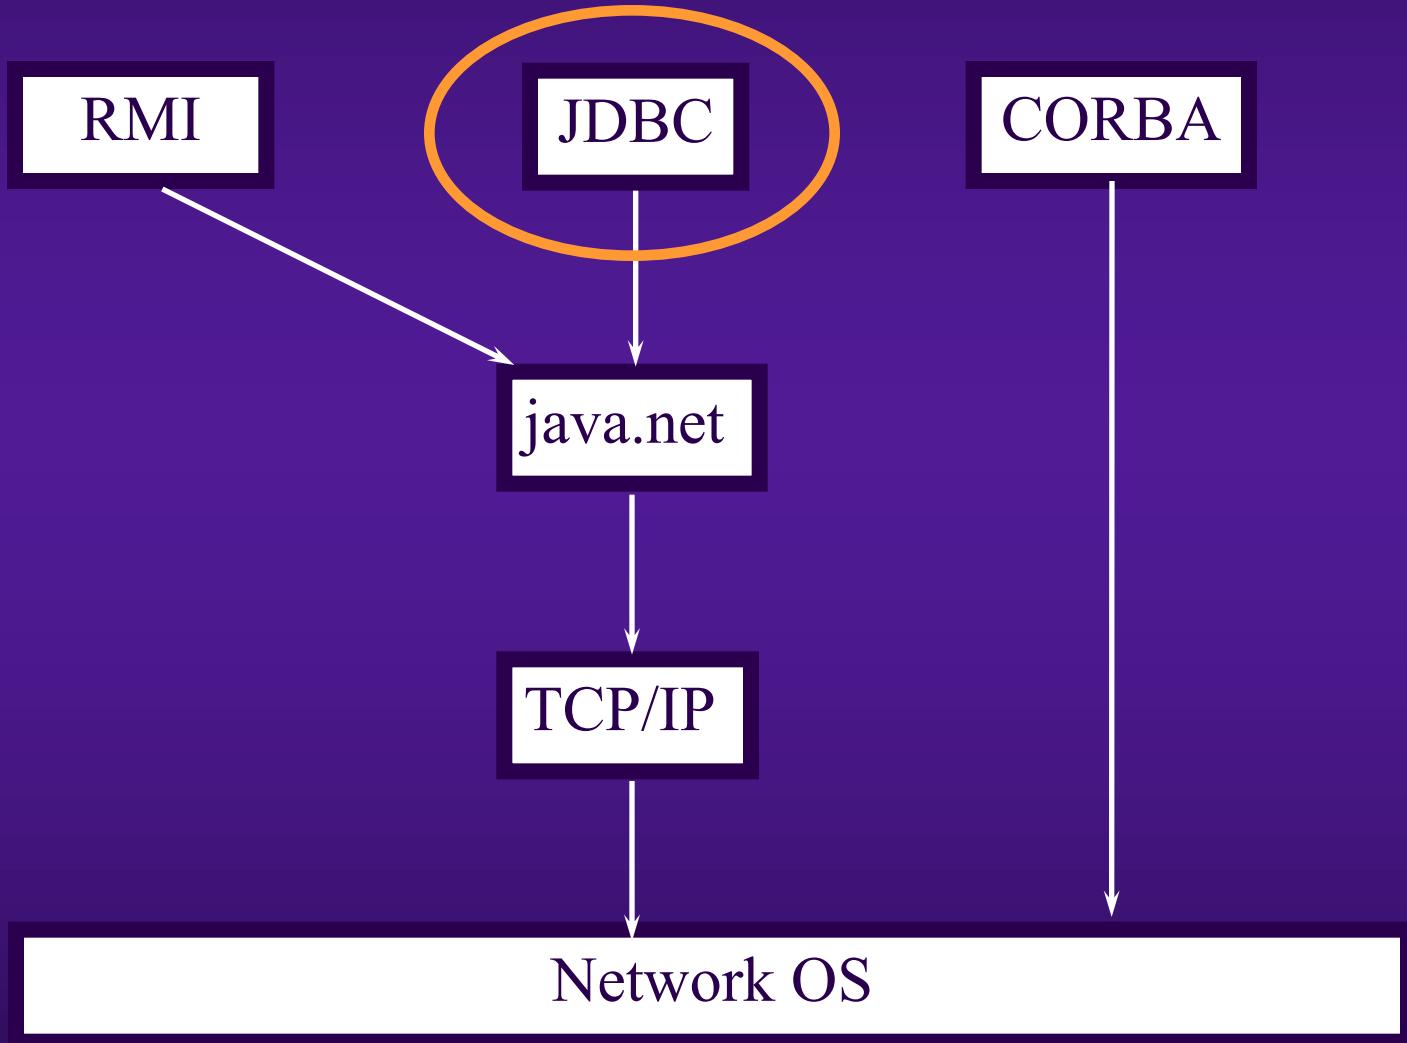


Agenda

- ◆ Overview of Databases and Java
- ◆ Overview of JDBC
- ◆ JDBC APIs
- ◆ Other Database Techniques



Overview





Why Java?

- ◆ Write once, run anywhere
 - ◆ Multiple client and server platforms
- ◆ Object-relational mapping
 - ◆ databases optimized for searching/indexing
 - ◆ objects optimized for engineering/flexibility
- ◆ Network independence
 - ◆ Works across Internet Protocol
- ◆ Database independence
 - ◆ Java can access any database vendor
- ◆ Ease of administration
 - ◆ zero-install client



Database Architectures

- ◆ Two-tier
- ◆ Three-tier
- ◆ N-tier



Database Technologies

- ◆ Hierarchical
 - ◆ obsolete (in a manner of speaking)
 - ◆ any specialized file format can be called a hierarchical DB
- ◆ Relational (aka SQL) (RDBMS)
 - ◆ row, column
 - ◆ most popular
- ◆ Object-relational DB (ORDBMS)
 - ◆ add inheritance, blobs to RDB
 - ◆ NOT object-oriented -- “object” is mostly a marketing term
- ◆ Object-oriented DB (OODB)
 - ◆ data stored as objects
 - ◆ high-performance for OO data models



Transactions

- ◆ Transaction = more than one statement which must all succeed (or all fail) together
- ◆ If one fails, the system must reverse all previous actions
- ◆ Also can't leave DB in inconsistent state halfway through a transaction
- ◆ COMMIT = complete transaction
- ◆ ROLLBACK = abort



JDBC Goals

- ◆ SQL-Level
- ◆ 100% Pure Java
- ◆ Keep it simple
- ◆ High-performance
- ◆ Leverage existing database technology
 - ◆ why reinvent the wheel?
- ◆ Use strong, static typing wherever possible
- ◆ Use multiple methods to express multiple functionality



JDBC Architecture



- ◆ Java code calls JDBC library
- ◆ JDBC loads a *driver*
- ◆ Driver talks to a particular database
- ◆ Can have more than one driver -> more than one database
- ◆ Ideal: can change database engines without changing any application code

General Architecture

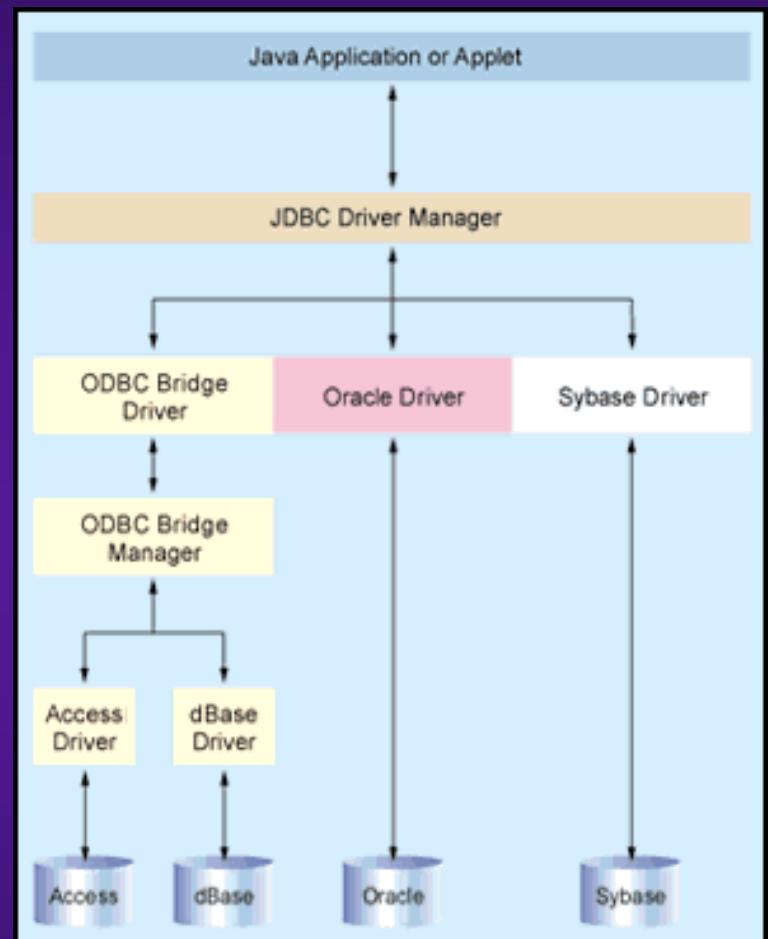
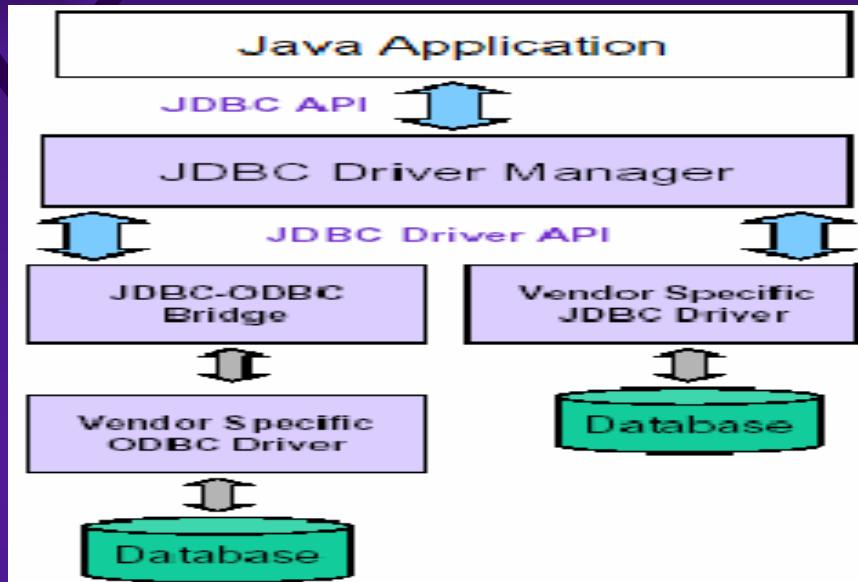


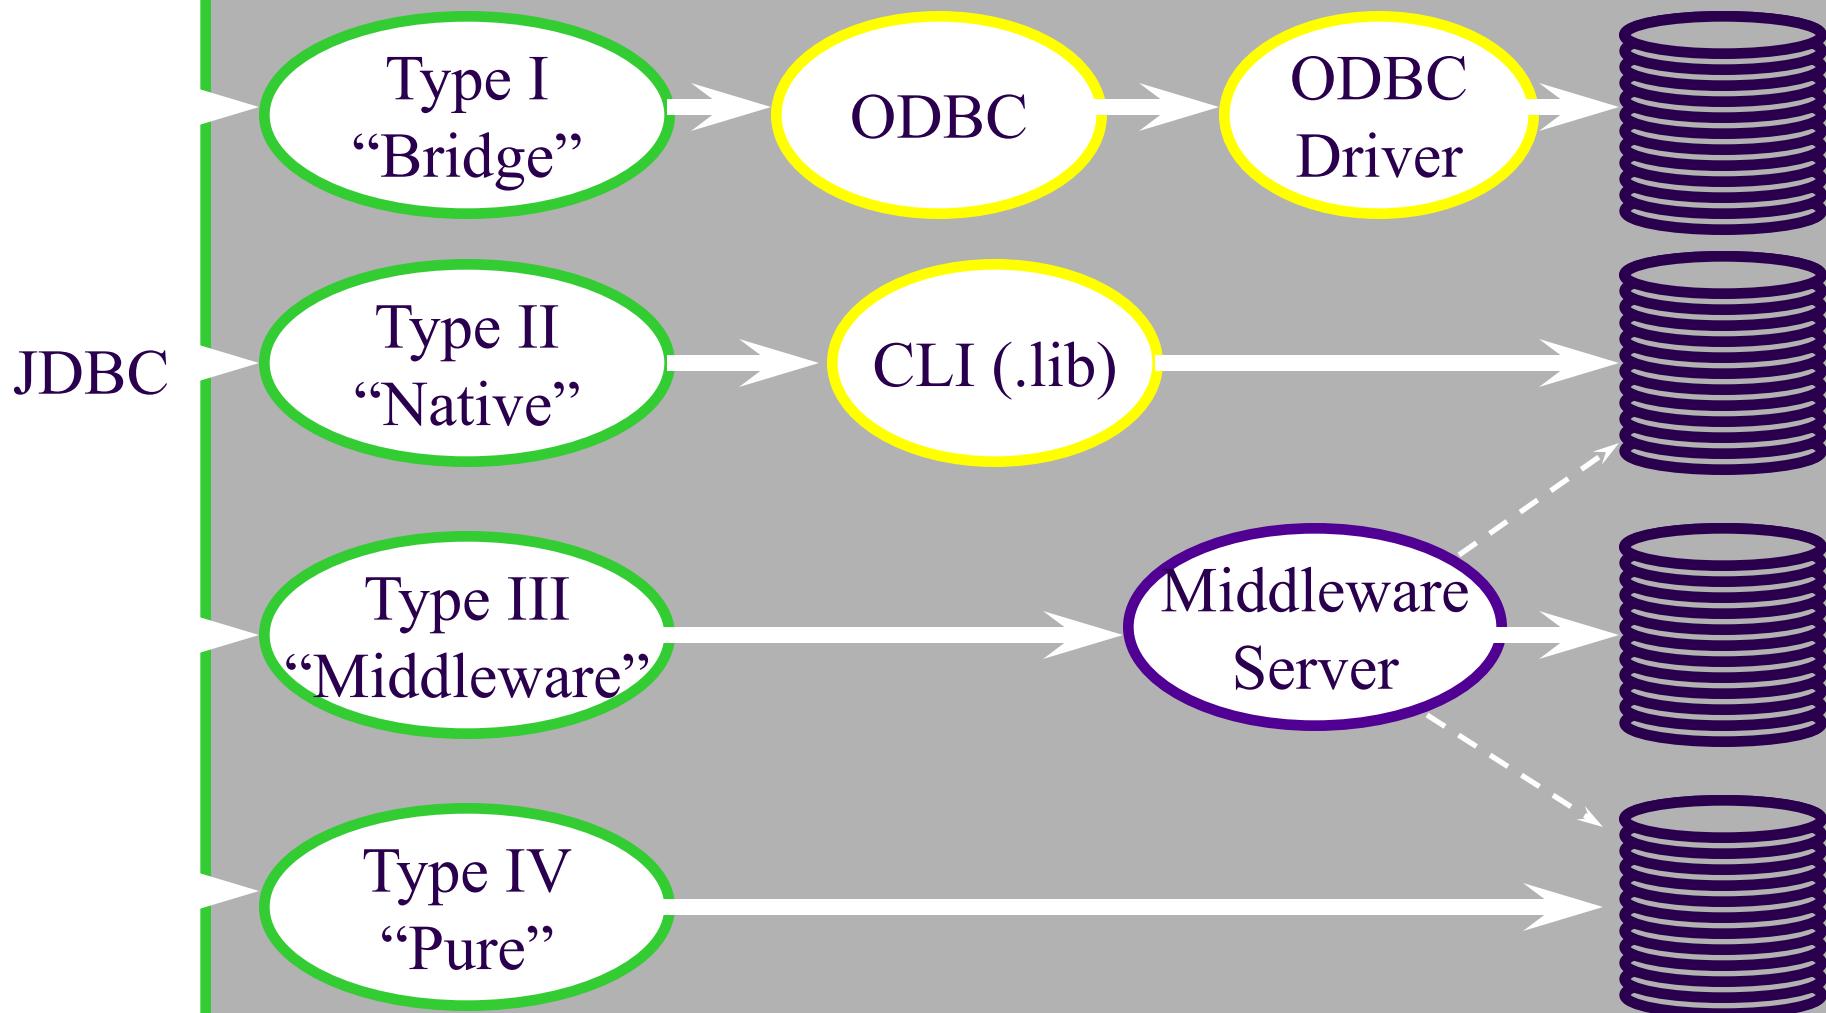
Figure 1. Anatomy of Data Access. The Driver Manager provides a consistent layer between your Java app and back-end database. JDBC works natively (such as with the Oracle driver in this example) or with any ODBC datasource.



JDBC Drivers

- ◆ Type I: “Bridge”
- ◆ Type II: “Native”
- ◆ Type III: “Middleware”
- ◆ Type IV: “Pure”

JDBC Drivers (Fig.)





JDBC Limitations

- ◆ No scrolling cursors
- ◆ No bookmarks



java.sql Package

Jdbc Classes

- **DriverManager**
 - Manages JDBC Drivers
 - Used to Obtain a connection to a Database
- **Types**
 - Defines constants which identify SQL types
- **Date**
 - Used to Map between java.util.Date and the SQL DATE type
- **Time**
 - Used to Map between java.util.Date and the SQL TIME type
- **TimeStamp**
 - Used to Map between java.util.Date and the SQL TIMESTAMP type



JDBC Interfaces

- **Driver**

- All JDBC Drivers must implement the Driver interface. Used to obtain a connection to a specific database type

- **Connection**

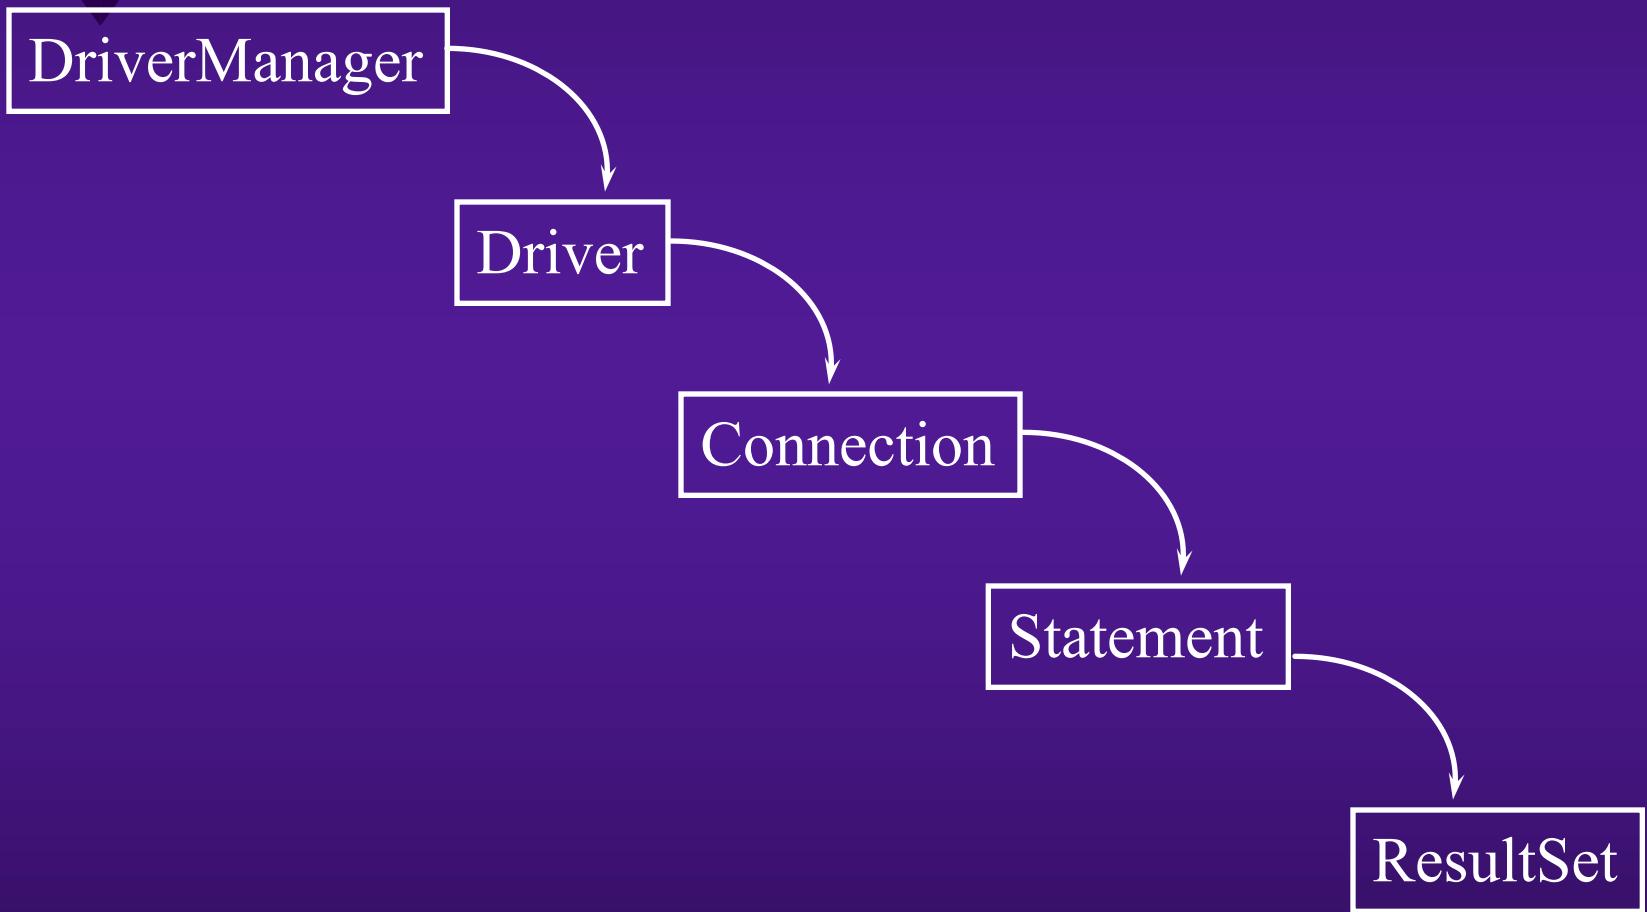
- Represents a connection to a specific database
- Used for creating statements
- Used for managing database transactions
- Used for accessing stored procedures
- Used for creating callable statements

- **Statement**

- Used for executing SQL statements against the database



JDBC Class Usage





Using JDBC

- To execute a statement against a database, the following flow is observed
 - Load the driver (Only performed once)
 - Obtain a Connection to the database (Save for later use)
 - Obtain a Statement object from the Connection
 - Use the Statement object to execute SQL. Updates, inserts and deletes return Boolean. Selects return a ResultSet
 - Navigate ResultSet, using data as required
 - Close ResultSet
 - Close Statement
- Do NOT close the connection
 - The same connection object can be used to create further statements
 - A Connection may only have one active Statement at a time. Do not forget to close the statement when it is no longer needed.
 - Close the connection when you no longer need to access the



Loading a Driver

- How does one load a "class" into the Virtual machine?
 - Use the static method Class.forName()
`Class.forName("sun.jdbc.odbc.JdbcOdbcDriver");`
- A connection is obtained in the following manner:

```
Connection aConnection =  
    DriverManager.getConnection("jdbc:odbc:myDatabase");
```

- Overloaded versions of the getConnection method allow the specification of a username and



MySQL	com.mysql.jdbc.Driver	jdbc:mysql://hostname/ databaseName
ORACLE	oracle.jdbc.driver.OracleDriver	jdbc:oracle:thin:@ hostname:port Number:databaseName
DB2	COM.ibm.db2.jdbc.net.DB2Driver	jdbc:db2: hostname:port Number/databaseName
Sybase	com.sybase.jdbc.SybDriver	jdbc:sybase:Tds: hostname: port Number/databaseName

```
String URL = "jdbc:oracle:thin:@amrood:1521:EMP"; String USER = "username"; String PASS  
= "password"
```



```
Connection getConnection  
    (String url, String user,  
     String password)
```

- ◆ Connects to given JDBC URL with given user name and password
- ◆ Throws java.sql.SQLException
- ◆ returns a Connection object



Connection Methods

Statement createStatement()

- ◆ returns a new Statement object

PreparedStatement prepareStatement(String sql)

- ◆ returns a new PreparedStatement object

CallableStatement prepareCall(String sql)

- ◆ returns a new CallableStatement object
- ◆ Why all these different kinds of statements?
 - ◆ Optimization.



Statement

- ◆ A Statement object is used for executing a static SQL statement and obtaining the results produced by it.



Statement Methods

`ResultSet executeQuery(String)`

- ◆ Execute a SQL statement that returns a single `ResultSet`.

`int executeUpdate(String)`

- ◆ Execute a SQL INSERT, UPDATE or DELETE statement. Returns the number of rows changed.

`boolean execute(String)`

- ◆ Execute a SQL statement that may return multiple results.
- ◆ Why all these different kinds of queries?
Optimization.



ResultSet

- ◆ A ResultSet provides access to a table of data generated by executing a Statement.
- ◆ Only one ResultSet per Statement can be open at once.
- ◆ The table rows are retrieved in sequence.
- ◆ A ResultSet maintains a cursor pointing to its current row of data.
- ◆ The 'next' method moves the cursor to the next row.
 - ◆ you can't rewind



ResultSet Methods

- ◆ boolean next()
 - ◆ activates the next row
 - ◆ the first call to next() activates the first row
 - ◆ returns false if there are no more rows
- ◆ void close()
 - ◆ disposes of the ResultSet
 - ◆ allows you to re-use the Statement that created it
 - ◆ automatically called by most Statement methods



ResultSet Methods

- ◆ *Type* `getType(int columnIndex)`
 - ◆ returns the given field as the given type
 - ◆ fields indexed starting at 1 (not 0)
- ◆ *Type* `getType(String columnName)`
 - ◆ same, but uses name of field
 - ◆ less efficient
- ◆ `int findColumn(String columnName)`
 - ◆ looks up column index given column name



ResultSet Methods

- ◆ `String getString(int columnIndex)`
- ◆ `boolean getBoolean(int columnIndex)`
- ◆ `byte getByte(int columnIndex)`
- ◆ `short getShort(int columnIndex)`
- ◆ `int getInt(int columnIndex)`
- ◆ `long getLong(int columnIndex)`
- ◆ `float getFloat(int columnIndex)`
- ◆ `double getDouble(int columnIndex)`
- ◆ `Date getDate(int columnIndex)`
- ◆ `Time getTime(int columnIndex)`
- ◆ `Timestamp getTimestamp(int columnIndex)`



SELECT Example

```
import java.sql.*;  
  
public class classname {  
    public static void main(String args[]) {  
        String url = "jdbc:odbc:dsn";  
        Connection con;  
        Statement st;  
        Class.forName("sun.jdbc.odbc.JdbcOdbcDriver");  
        con = DriverManager.getConnection(url,  
            "alex", "8675309");  
        st = con.createStatement();  
        ResultSet results = st.executeQuery("SELECT  
            EmployeeID, LastName, FirstName FROM  
            Employees");
```



SELECT Example (Cont.)

```
while (results.next()) {  
    int id = results.getInt(1);  
    String last = results.getString(2);  
    String first = results.getString(3);  
    System.out.println(" " + id + ": " +  
        first + " " + last);  
}  
st.close();  
con.close();  
}  
}
```



Mapping Java Types to SQL Types

<u>SQL type</u>	<u>Java Type</u>
CHAR, <u>VARCHAR</u> , LONGVARCHAR	String
<u>NUMERIC</u> , DECIMAL	java.math.BigDecimal
BIT	boolean
TINYINT	byte
SMALLINT	short
INTEGER	int
BIGINT	long
REAL	float
FLOAT, <u>DOUBLE</u>	double
BINARY, <u>VARBINARY</u> , LONGVARBINARY	byte[]
DATE	java.sql.Date
TIME	java.sql.Time
TIMESTAMP	java.sql.Timestamp



Executing Dynamic SQL Queries

```
import java.sql.*;  
public class JdbcPreparedStatementExample {  
    static private final String driver =  
    "sun.jdbc.odbc.JdbcOdbcDriver";  
    static private final String connection = "jdbc:odbc:emp";  
    public static void main(String args[]) {  
        Connection con = null;  
        PreparedStatement pst = null;  
        ResultSet rs = null;  
        try {  
            Class.forName(driver);  
            con = DriverManager.getConnection(connection);  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
        String sql = "SELECT * FROM emp WHERE id = ?";  
        try {  
            pst = con.prepareStatement(sql);  
            pst.setInt(1, 100);  
            rs = pst.executeQuery();  
            while (rs.next()) {  
                System.out.println(rs.getString("name") + " " +  
                    rs.getInt("id") + " " +  
                    rs.getDouble("sal"));  
            }  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```



```
String sql = "select * from Employees where FirstName " + "in(?,?,?)";
pst = con.prepareStatement(sql);
pst.setString(1, "komal");
pst.setString(2, "ajay");
pst.setString(3, "santosh");
rs = pst.executeQuery();
while (rs.next()) {
    System.out.println(rs.getString(1));
    System.out.println(rs.getString(2));
    System.out.println(rs.getString(3));
}
} catch (Exception e) {
    System.out.println(e);
}
}}
```



Executing Stored Procedures

```
CallableStatement cs2 = con.prepareCall("{call  
ADDITION(?, ?, ?)}");  
cs2.registerOutParameter(3,java.sql.Types.INTEGER);  
    cs2.setInt(1,10);  
    cs2.setInt(2,25);  
    cs2.execute();  
    int res = cs2.getInt(3);  
    System.out.println(res);
```

```
CREATE OR REPLACE PROCEDURE getEmpName  
(EMP_ID IN NUMBER, EMP_FIRST OUT VARCHAR)  
AS BEGIN SELECT first INTO EMP_FIRST FROM  
Employees WHERE ID = EMP_ID; END;
```



Modifying the Database

- ◆ use executeUpdate if the SQL contains “INSERT” or “UPDATE”
- ◆ Why isn’t it smart enough to parse the SQL?
Optimization.
- ◆ executeUpdate returns the number of rows modified
- ◆ executeUpdate also used for “CREATE TABLE” etc. (DDL)



Transaction Management

- ◆ The **atomicity** of a transaction means that if any part of it fails, the entire transaction is aborted. So the transaction is committed only if each part of it executes successfully.
- ◆ The **isolation** property ensures that transactions don't interfere with each other's processing.
- ◆ The **durability** of transactions ensures that all changes made to the databases by a transaction are permanent once that transaction is committed.
- ◆ The **consistency** of a transaction ensures that when it is committed, all relevant databases are in a consistent state.



Transaction Management

- ◆ Transactions are not explicitly opened and closed
- ◆ Instead, the connection has a state called *AutoCommit* mode
- ◆ if *AutoCommit* is true, then every statement is automatically committed
- ◆ default case: true



setAutoCommit

`Connection.setAutoCommit(boolean)`

- ◆ if *AutoCommit* is false, then every statement is added to an ongoing transaction
- ◆ you must explicitly commit or rollback the transaction using `Connection.commit()` and `Connection.rollback()`



- ◆ try{ //Assume a valid connection object conn
conn.setAutoCommit(false);

Statement stmt = conn.createStatement();

String SQL = "INSERT INTO Employees " + "VALUES
(106, 20, 'Rita', 'Tez')";

stmt.executeUpdate(SQL);

//Submit a malformed SQL statement that breaks String SQL
= "INSERTED IN Employees " + "VALUES (107, 22, 'Sita',
'Singh')";

stmt.executeUpdate(SQL); // If there is no error.
conn.commit();
}
catch(SQLException se)
{ // If there is any error. conn.rollback(); }



Batch Processing allows to group related SQL statements into a batch and submit them with one call to the database.

- ◆ The **addBatch()** method of *Statement*, *PreparedStatement*, and *CallableStatement* is used to add individual statements to the batch. The **executeBatch()** is used to start the execution of all the statements grouped together.
- ◆ The **executeBatch()** returns an array of integers, and each element of the array represents the update count for the respective update statement.
- ◆ Just as you can add statements to a batch for processing, you can remove them with the **clearBatch()** method
- ◆

```
Statement stmt = conn.createStatement();
// Set auto-commit to false conn.setAutoCommit(false);
// Create SQL statement String SQL = "INSERT INTO Employees (id, first, last, age) "
+ "VALUES(200,'Zia', 'Ali', 30)";
// Add above SQL statement in the batch. stmt.addBatch(SQL);
// Create one more SQL statement String SQL = "INSERT INTO Employees (id, first,
last, age) " + "VALUES(201,'Raj', 'Kumar', 35)";
// Add above SQL statement in the batch. stmt.addBatch(SQL);
// Create one more SQL statement String SQL = "UPDATE Employees SET age = 35 "
+ "WHERE id = 100";
// Add above SQL statement in the batch. stmt.addBatch(SQL);
// Create an int[] to hold returned values int[] count = stmt.executeBatch();
//Explicitly commit statements to apply changes conn.commit();
```



Javax package

- ◆ Javax.sql.*Automatic driver loading
- ◆ Exception handling improvements
- ◆ Enhanced BLOB/CLOB functionality
- ◆ Connection and statement interface enhancements
- ◆ XML datatype support



- Connection Pooling is a **Server-side** technology
- It is extremely expensive to open and close database connections every time a user requests a connection to carry out a SQL statement.
- With connection pooling, the application server maintains a pre-defined pool of “open” database connections that are shared between application clients
- When the client creates a connection via a **DataSource** object, the connection will come out of the pool.
- When a client closes a connection, the connection is returned back to the pool.

- 
- ◆

```
import java.sql.Connection;
import javax.naming.Context;
import javax.naming.InitialContext;
import javax.sql.DataSource;

Connection conn = null;

try {
    Context      ctx      = new InitialContext();
    DataSource   ds       = (DataSource) ctx.lookup("jdbc/MurachBooks");
    Connection   conn     = ds.getConnection( "myLogin", "myPassword");
}
catch (Exception e) { e.printStackTrace(); }

try {
    PreparedStatement ps = conn.prepareStatement( "select * from file" );
    ResultSet rs = ps.executeQuery();
    while (rs.next()) { ... }
}
```
 - ◆ `javax.sql.RowSet`



JDBC Class Diagram

